



BUILDING RECOMMENDATION ENGINES IN PYSPARK

Introduction to the MovieLens dataset

Jamen Long
Data Scientist



MOVIELENS DATASET:

F. Maxwell Harper and Joseph A. Konstan. 2015

The MovieLens Datasets: History and Context.

ACM Transactions on Interactive Intelligent Systems (TiiS) 5, 4, Article 19 (December 2015), 19 Pages. DOI=<http://dx.doi.org/10.1145/2827872>



MOVIELENS DATASET:

F. Maxwell Harper and Joseph A. Konstan. 2015
The MovieLens Datasets: History and Context.
ACM Transactions on Interactive Intelligent Systems (TiiS) 5, 4, Article 19 (December 2015), 19 Pages. DOI=<http://dx.doi.org/10.1145/2827872>

Ratings: 20,00263

Users: 138,493

Movies: 27,278



Explore the Data

```
df.show()  
df.columns()
```



MovieLens Sparsity

$$Sparsity = \frac{\text{Number of Ratings in Matrix}}{(\text{Number of Users}) \times (\text{Number of Movies})}$$



Sparsity: Numerator

```
# Number of ratings in matrix  
numerator = ratings.count()
```



Sparsity: Users and Movies

```
# Distinct users and movies
users = ratings.select("userId").distinct().count()
movies = ratings.select("movieId").distinct().count()
```

Sparsity: Denominator

```
# Number of ratings in matrix
numerator = ratings.count()

# Distinct users and movies
users = ratings.select("userId").distinct().count()
movies = ratings.select("movieId").distinct().count()

# Number of ratings matrix could contain if no empty cells
denominator = users * movies
```




Sparsity

```
# Number of ratings in matrix
numerator = ratings.count()

# Distinct users and movies
users = ratings.select("userId").distinct().count()
movies = ratings.select("movieId").distinct().count()

# Number of ratings matrix could contain if no empty cells
denominator = users * movies

#Calculating sparsity
sparsity = 1 - (numerator*1.0 / denominator)
print ("Sparsity: "), sparsity
```

```
Sparsity: .998
```



The .distinct() Method

```
ratings.select("userId").distinct().count()
```

```
671
```



GroupBy Method

```
# Group by userId  
ratings.groupBy("userId")
```

GroupBy Method

```
# Num of song plays by userId
ratings.groupBy("userId").count().show()
```

```
+-----+-----+
|userId|count|
+-----+-----+
|    148|    76|
|    243|    12|
|     31|   232|
|    137|    16|
|    251|    19|
|     85|   752|
|     65|   737|
|    255|     9|
|     53|   190|
|    133|   302|
|    296|    74|
|     78|   301|
|    108|   136|
|    155|     3|
|    193|   174|
|    101|     1|
+-----+-----+
```

GroupBy Method Min

```
from pyspark.sql.functions import min, max, avg

# Min num of song plays by userId
msd.groupBy("userId").count()
      .select(min("count")).show()
```

```
+-----+
|min(count)|
+-----+
|          1|
+-----+
```

GroupBy Method Max

```
from pyspark.sql.functions import min, max, avg

# Min num of song plays by userId
ratings.groupBy("userId").count()
        .select(min("count")).show()
```

```
+-----+
|min(count)|
+-----+
|          56|
+-----+
```

```
# Max num of song plays by userId
ratings.groupBy("userId").count()
        .select(max("count")).show()
```

```
+-----+
|max(count)|
+-----+
|         1162|
+-----+
```



GroupBy Method Avg

```
# Avg num of song plays by userId
ratings.groupBy("userId").count()
        .select(avg("count")).show()
```

```
+-----+
|avg(count)|
+-----+
| 233.34579|
+-----+
```



Filter Method

```
# Removes users with less than 20 ratings
ratings.groupBy("userId").count().filter(col("count") >= 20).show()
```

```
+-----+-----+
|userId|count|
+-----+-----+
|    148|    76|
|     31|   232|
|     85|   752|
|     65|   737|
|     53|   190|
|    133|   302|
|    296|    74|
|     78|   301|
|    108|   136|
|    193|   174|
+-----+-----+
```




BUILDING RECOMMENDATION ENGINES IN PYSPARK

Let's practice!



BUILDING RECOMMENDATION ENGINES IN PYSPARK

ALS model buildout on MovieLens Data

Jamen Long
Data Scientist

Fitting a Basic Model

```
# Split data
(training_data, test_data) = movie_ratings.randomSplit([0.8, 0.2])

# Build ALS model
from pyspark.ml.recommendation import ALS

als = ALS(userCol="userId", itemCol="movieId", ratingCol="rating",
          rank=25, maxIter=100, regParam=.05, nonnegative=True,
          coldStartStrategy="drop", implicitPrefs=False)

# Fit model to training data
model = als.fit(training_data)

# Generate predictions on test_data
predictions = model.transform(test_data)

# Tell Spark how to evaluate predictions
evaluator = RegressionEvaluator(metricName="rmse", labelCol="rating",
                               predictionCol="prediction")

# Obtain and print RMSE
rmse = evaluator.evaluate(predictions)
print ("RMSE: "), rmse
```

```
RMSE: 1.45
```



Intro to ParamGridBuilder and CrossValidator

```
ParamGridBuilder()
```

```
CrossValidator()
```



ParamGridBuilder

```
# Imports ParamGridBuilder package
from pyspark.ml.tuning import ParamGridBuilder

# Creates a ParamGridBuilder
param_grid = ParamGridBuilder()
```



Adding Hyperparameters to the ParamGridBuilder

```
# Imports ParamGridBuilder package
from pyspark.ml.tuning import ParamGridBuilder

# Creates a ParamGridBuilder, and adds hyperparameters
param_grid = ParamGridBuilder()
                .addGrid(als.rank, [])
                .addGrid(als.maxIter, [])
                .addGrid(als.regParam, [])
```



Adding Hyperparameter Values to the ParamGridBuilder

```
# Imports ParamGridBuilder package
from pyspark.ml.tuning import ParamGridBuilder

# Creates a ParamGridBuilder, and adds hyperparameters and values
param_grid = ParamGridBuilder()
    .addGrid(als.rank, [5, 40, 80, 120])
    .addGrid(als.maxIter, [5, 100, 250, 500])
    .addGrid(als.regParam, [.05, .1, 1.5])
    .build()
```



CrossValidator

```
# Imports CrossValidator package
from pyspark.ml.tuning import CrossValidator

# Creates cross validator and tells Spark what to use when training # and evaluation
cv = CrossValidator(estimator = als,
                    estimatorParamMaps = param_grid,
                    evaluator = evaluator,
                    numFolds = 5)
```




Cross Validator Instantiation and Estimator

```
# Imports CrossValidator package
from pyspark.ml.tuning import CrossValidator

# Instantiates a cross validator
cv = CrossValidator()
```



Cross Validator ParamMaps

```
# Imports CrossValidator package
from pyspark.ml.tuning import CrossValidator

# Tells Spark what to use when training a model
cv = CrossValidator(estimator = als,
                    estimatorParamMaps = param_grid,
                    )
```



Cross Validator

```
# Imports CrossValidator package
from pyspark.ml.tuning import CrossValidator

# Tells Spark what alg, hyperparameter values, how to evaluate
# each model and number of folds to use during training
cv = CrossValidator(estimator = als,
                    estimatorParamMaps = param_grid,
                    evaluator = evaluator,
                    numFolds = 5)
```



Random Split

```
# Create training and test set (80/20 split)
(training, test) = movie_ratings.randomSplit([0.8, 0.2])

# Build generic ALS model without hyperparameters
als = ALS(userCol="userId", itemCol="movieId", ratingCol="rating",
          coldStartStrategy="drop", nonnegative = True,
          implicitPrefs = False)
```



ParamGridBuilder

```
# Create training and test set (80/20 split)
(training, test) = movie_ratings.randomSplit([0.8, 0.2])

# Build generic ALS model without hyperparameters
als = ALS(userCol="userId", itemCol="movieId", ratingCol="rating",
          coldStartStrategy="drop", nonnegative = True,
          implicitPrefs = False)

# Tell Spark what values to try for each hyperparameter
from pyspark.ml.tuning import ParamGridBuilder
param_grid = ParamGridBuilder()
    .addGrid(als.rank, [5, 40, 80, 120])
    .addGrid(als.maxIter, [5, 100, 250, 500])
    .addGrid(als.regParam, [.05, .1, 1.5])
    .build()
```


Cross Validator

```
# Build generic ALS model without hyperparameters
als = ALS(userCol="userId", itemCol="movieId", ratingCol="rating",
          coldStartStrategy="drop", nonnegative = True,
          implicitPrefs = False)

# Tell Spark what values to try for each hyperparameter
from pyspark.ml.tuning import ParamGridBuilder
param_grid = ParamGridBuilder()
    .addGrid(als.rank, [5, 40, 80, 120])
    .addGrid(als.maxIter, [5, 100, 250, 500])
    .addGrid(als.regParam, [.05, .1, 1.5])
    .build()

# Tell Spark how to evaluate model performance
evaluator = RegressionEvaluator(metricName="rmse", labelCol="rating",
                               predictionCol="prediction")

# Build cross validation step using CrossValidator
from pyspark.ml.tuning import CrossValidator
cv = CrossValidator(estimator = als,
                   estimatorParamMaps = param_grid,
                   evaluator = evaluator,
                   numFolds = 5)
```



Best Model

```
# Tell Spark what values to try for each hyperparameter
from pyspark.ml.tuning import ParamGridBuilder
param_grid = ParamGridBuilder()
    .addGrid(als.rank, [5, 40, 80, 120])
    .addGrid(als.maxIter, [5, 100, 250, 500])
    .addGrid(als.regParam, [.05, .1, 1.5])
    .build()

# Tell Spark how to evaluate model performance
evaluator = RegressionEvaluator(metricName="rmse", labelCol="rating",
                                predictionCol="prediction")

# Build cross validation step using CrossValidator
from pyspark.ml.tuning import CrossValidator
cv = CrossValidator(estimator = als,
                    estimatorParamMaps = param_grid,
                    evaluator = evaluator,
                    numFolds = 5)

# Run the cv on the training data
model = cv.fit(training)

# Extract best combination of values from cross validation
best_model = model.bestModel
```


Predictions and Performance Evaluation

```
# Extract best combination of values from cross validation
best_model = model.bestModel

# Generate test set predictions and evaluate using RMSE
predictions = best_model.transform(test)
rmse = evaluator.evaluate(predictions)

# Print evaluation metrics and model parameters
print (**Best Model**)
print ("RMSE = "), rmse
print (" Rank: "), best_model.rank
print (" MaxIter: "), best_model._java_obj.parent().getMaxIter()
print (" RegParam: "), best_model._java_obj.parent().getRegParam()
```



BUILDING RECOMMENDATION ENGINES IN PYSPARK

Let's practice!



BUILDING RECOMMENDATION ENGINES IN PYSPARK

Model Performance Evaluation and Output Cleanup

Jamen Long
Data Scientist



Root Mean Squared Error

$$\text{RMSE} = \sqrt{\frac{\sum (y_{\text{pred}} - y_{\text{actual}})^2}{N}}$$



Pred vs Actual

```
+-----+-----+
|pred|actual|
+-----+-----+
|    5|    4.5|
|    3|    3.5|
|    4|     4|
|    2|     1|
+-----+-----+
```



Pred vs Actual: Difference

```
+-----+-----+-----+
|pred|actual|diff|
+-----+-----+-----+
|    5|    4.5| 0.5|
|    3|    3.5|-0.5|
|    4|     4| 0.0|
|    2|     1| 1.0|
+-----+-----+-----+
```



Difference Squared

```
+-----+-----+-----+-----+
|pred|actual|diff|diff_sq|
+-----+-----+-----+-----+
|  5 |   4.5 | 0.5 |   0.25 |
|  3 |   3.5 |-0.5 |   0.25 |
|  4 |    4 | 0.0 |   0.00 |
|  2 |    1 | 1.0 |   1.00 |
+-----+-----+-----+-----+
```



Sum of Difference Squared

```
+-----+-----+-----+-----+
|pred|actual|diff|diff_sq|
+-----+-----+-----+-----+
|  5 |   4.5 | 0.5 |   0.25 |
|  3 |   3.5 |-0.5 |   0.25 |
|  4 |    4 | 0.0 |   0.00 |
|  2 |    1 | 1.0 |   1.00 |
+-----+-----+-----+-----+
```

```
sum of diff_sq = 1.5
```




Average of Difference Squared

```
+-----+-----+-----+-----+
|pred|actual|diff|diff_sq|
+-----+-----+-----+-----+
|  5|  4.5| 0.5|  0.25|
|  3|  3.5|-0.5|  0.25|
|  4|  4| 0.0|  0.00|
|  2|  1| 1.0|  1.00|
+-----+-----+-----+-----+
```

sum of diff_sq = 1.5

avg of diff_sq = 1.5 / 4 = 0.375



RMSE

```
+-----+-----+-----+-----+
|pred|actual|diff|diff_sq|
+-----+-----+-----+-----+
|  5 |   4.5 | 0.5 |   0.25 |
|  3 |   3.5 |-0.5 |   0.25 |
|  4 |    4 | 0.0 |   0.00 |
|  2 |    1 | 1.0 |   1.00 |
+-----+-----+-----+-----+
```

```
sum of diff_sq = 1.5
avg of diff_sq = 1.5 / 4 = 0.375
RMSE = sq root of avg of diff_sq = 0.61
```



Recommend for all users

```
# Generate n recommendations for all users  
recommendForAllUsers(n) # n is an integer
```



Unclean Recommendation Output

```
ALS_recommendations.show()
```

```
+-----+-----+
|userId| recommendations|
+-----+-----+
|    360| [[65037, 4.491346] ...|
|    246| [[3414, 4.8967672] ...|
|    346| [[4565, 4.9247236] ...|
|    476| [[83318, 4.9556283] ...|
|    367| [[4632, 4.7018986] ...|
|    539| [[1172, 5.2528191] ...|
|    599| [[6413, 4.7284415] ...|
|    220| [[80, 4.4857406] ...|
|    301| [[66665, 5.190159] ...|
|    173| [[65037, 4.316745] ...|
+-----+-----+
```



Cleaning Up Recommendation Output

```
ALS_recommendations.registerTempTable("ALS_recs_temp")

clean_recs = spark.sql("""SELECT userId,
                           movieIds_and_ratings.movieId AS movieId,
                           movieIds_and_ratings.rating AS prediction
                           FROM ALS_recs_temp
                           LATERAL VIEW explode(recommendations) exploded_table
                           AS movieIds_and_ratings""")
```

Explode Function

```
exploded_recs = spark.sql("SELECT uderId,  
                             explode(recommendations) AS MovieRec  
                             FROM ALS_recs_temp")  
  
exploded_recs.show()
```

```
+-----+-----+  
|userId|                               MovieRec|  
+-----+-----+  
|   360|{"movieId": 65037, "rating": 4.4913464}|  
|   360|{"movieId": 59684, "rating": 4.4832921}|  
|   360|{"movieId": 31435, "rating": 4.4822811}|  
|   360|{"movieId": 593, "rating": 4.456215}|  
|   360|{"movieId": 67504, "rating": 4.4028492}|  
|   360|{"movieId": 83411, "rating": 4.3391834}|  
|   360|{"movieId": 83318, "rating": 4.3199939}|  
|   360|{"movieId": 83359, "rating": 4.3000213}|  
|   360|{"movieId": 76170, "rating": 4.2987138}|  
|   360|{"movieId": 17, "rating": 4.2539403}|  
|   360|{"movieId": 2112, "rating": 4.11893843}|  
+-----+-----+
```



Adding Lateral View

```
ALS_recommendations.registerTempTable("ALS_recs_temp")

clean_recs = spark.sql("SELECT userId,
                        movieIds_and_ratings.movieId AS movieId,
                        movieIds_and_ratings.rating AS prediction
                        FROM ALS_recs_temp
                        LATERAL VIEW explode(recommendations) exploded_table
                        AS movieIds_and_ratings")
```

Explode and Lateral View Together

```
ALS_recommendations.registerTempTable("ALS_recs_temp")

clean_recs = spark.sql("""SELECT userId,
                           movieIds_and_ratings.movieId AS movieId,
                           movieIds_and_ratings.rating AS prediction
                           FROM ALS_recs_temp
                           LATERAL VIEW explode(recommendations) exploded_table
                           AS movieIds_and_ratings""")

clean_recs.show()
```

```
+-----+-----+
|userId|movieId|prediction|
+-----+-----+
|    360|   65037|   4.491346|
|    360|   59684|   4.491346|
|    360|   34135|   4.491346|
|    360|     593|   4.453185|
|    360|   67504|   4.389951|
|    360|   83411|   4.389944|
|    360|   83318|   4.389938|
|    360|   83359|   4.373281|
|    360|   76173|   4.190159|
|    360|    5114|   4.116745|
+-----+-----+
```




```
clean_recs.join(movie_info, ["movieId"], "left").show()
```

```
+-----+-----+-----+-----+
|userId|movieId|prediction|          title|
+-----+-----+-----+-----+
|    360|   65037|   4.491346|    Ben X (2007)|
|    360|   59684|   4.491346|  Lake of Fire (2006)|
|    360|   34135|   4.491346| Rory O Shea Was H...|
|    360|     593|   4.453185| Silence of the La...|
|    360|   67504|   4.389951| Land of Silence a...|
|    360|   83411|   4.389944|         Cops (1922)|
|    360|   83318|   4.389938|    Goat, The (1921)|
|    360|   83359|   4.373281| Play House, The(...|
|    360|   76173|   4.190159| Micmacs (Micmacs...|
|    360|    5114|   4.116745| Bad and the Beaut...|
+-----+-----+-----+-----+
```



Filtering Recommendations

```
clean_recs.join(movie_ratings, ["userId", "movieId"], "left")
```



```
clean_recs.join(movie_ratings, ["userId", "movieId"], "left").show()
```

```
+-----+-----+-----+
|userId|movieId|prediction|rating|
+-----+-----+-----+
|    173|    318|  4.947126|  null|
|    150|    318|  4.066513|   5.0|
|    369|    318|  4.514297|   5.0|
|     27|    318|  4.523860|  null|
|     42|    318|  4.568357|   5.0|
|    662|    318|  4.242076|   5.0|
|    250|    318|  5.042126|   5.0|
|     94|    318|  4.291757|   5.0|
|    515|    318|  5.165822|  null|
|    109|    318|  4.885314|   5.0|
+-----+-----+-----+
```



```
clean_recs.join(movie_ratings, ["userId", "movieId"], "left")  
            .filter(movie_ratings.rating.isNull()) .show()
```

```
+-----+-----+-----+  
|userId|movieId|prediction|rating|  
+-----+-----+-----+  
|    173|    318|  4.947126|  null|  
|     27|    318|  4.523860|  null|  
|    515|    318|  5.165822|  null|  
|    275|    318|  5.171431|  null|  
|    503|    318|  4.308533|  null|  
|    106|    318|  4.688634|  null|  
|    249|    318|  4.759836|  null|  
|    368|    318|  3.589334|  null|  
|    581|    318|  4.717382|  null|  
|    208|    318|  3.920525|  null|  
+-----+-----+-----+
```



BUILDING RECOMMENDATION ENGINES IN PYSPARK

Let's practice!