

```
In [1]: from pyspark.sql import SparkSession
import pyspark.sql.functions as F
from pyspark.sql.types import *

spark = SparkSession\
    .builder\
    .appName("chapter-24-ML")\
    .getOrCreate()

import os
SPARK_BOOK_DATA_PATH = os.environ['SPARK_BOOK_DATA_PATH']
```

```
In [2]: from pyspark.ml.linalg import Vectors
denseVec = Vectors.dense(1.0, 2.0, 3.0)
```

```
In [3]: denseVec
```

```
Out[3]: DenseVector([1.0, 2.0, 3.0])
```

```
In [7]: denseVec.array
```

```
Out[7]: array([1., 2., 3.])
```

```
In [8]: denseVec.values
```

```
Out[8]: array([1., 2., 3.])
```

```
In [9]: size = 3
idx = [1, 2] # locations of non-zero elements in vector
values = [2.0, 3.0]
sparseVec = Vectors.sparse(size, idx, values)
```

```
In [10]: sparseVec
```

```
Out[10]: SparseVector(3, {1: 2.0, 2: 3.0})
```

```
In [11]: sparseVec.values
```

```
Out[11]: array([2., 3.])
```

```
In [12]: # COMMAND -----
```

```
df = spark.read.json(SPARK_BOOK_DATA_PATH + "/data/simple-ml")
```

```
In [14]: df.count()
```

```
Out[14]: 110
```

```
In [15]: df.printSchema()
```

```
root
 |-- color: string (nullable = true)
 |-- lab: string (nullable = true)
 |-- value1: long (nullable = true)
 |-- value2: double (nullable = true)
```

```
In [13]: df.show(3)
```

```
+-----+-----+-----+-----+
|color| lab|value1|          value2|
+-----+-----+-----+-----+
|green|good|    1|14.386294994851129|
| blue|bad|    8|14.386294994851129|
| blue|bad|   12|14.386294994851129|
+-----+-----+-----+-----+
only showing top 3 rows
```

```
In [16]: df.orderBy("value1").show(10)
```

```
+-----+-----+-----+-----+
|color| lab|value1|          value2|
+-----+-----+-----+-----+
|green|good|      1|14.386294994851129|
|green|good|      1|14.386294994851129|
|  red| bad|      1| 38.97187133755819|
|green|good|      1|14.386294994851129|
|  red| bad|      1| 38.97187133755819|
|  red| bad|      1| 38.97187133755819|
|  red| bad|      1| 38.97187133755819|
|green|good|      1|14.386294994851129|
|  red| bad|      1| 38.97187133755819|
|green|good|      1|14.386294994851129|
+-----+-----+-----+-----+
only showing top 10 rows
```

```
In [18]: df.groupBy("color", "lab").count()\
         .orderBy("color", "lab")\
         .show(10)
```

```
+-----+-----+-----+
|color| lab|count|
+-----+-----+-----+
| blue| bad|    20|
|green| bad|    10|
|green|good|    30|
|  red| bad|    30|
|  red|good|    20|
+-----+-----+-----+
```

```
In [20]: # COMMAND -----
```

```
from pyspark.ml.feature import RFormula
supervised = RFormula(formula="lab ~ . + color:value1 + color:value2")
```

In [23]: `# COMMAND -----`

`## prepare feature columns`

```
fittedRF = supervised.fit(df)
preparedDF = fittedRF.transform(df)
preparedDF.show(10, False)
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+
|color|lab |value1|value2          | features
|label|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+
|green|good|1      | 14.386294994851129| (10, [1,2,3,5,8], [1.0,1.0,14.386294994851129,1.0,14.3862949948
51129]) |1.0 |
|blue |bad |8      | 14.386294994851129| (10, [2,3,6,9], [8.0,14.386294994851129,8.0,14.38629499485112
9]) |0.0 |
|blue |bad |12     | 14.386294994851129| (10, [2,3,6,9], [12.0,14.386294994851129,12.0,14.38629499485112
9]) |0.0 |
|green|good|15     | 38.97187133755819 | (10, [1,2,3,5,8], [1.0,15.0,38.97187133755819,15.0,38.971871337
55819]) |1.0 |
|green|good|12     | 14.386294994851129| (10, [1,2,3,5,8], [1.0,12.0,14.386294994851129,12.0,14.38629499
4851129]) |1.0 |
|green|bad |16     | 14.386294994851129| (10, [1,2,3,5,8], [1.0,16.0,14.386294994851129,16.0,14.38629499
4851129]) |0.0 |
|red  |good|35     | 14.386294994851129| (10, [0,2,3,4,7], [1.0,35.0,14.386294994851129,35.0,14.38629499
4851129]) |1.0 |
|red  |bad |1      | 38.97187133755819 | (10, [0,2,3,4,7], [1.0,1.0,38.97187133755819,1.0,38.97187133755
819]) |0.0 |
|red  |bad |2      | 14.386294994851129| (10, [0,2,3,4,7], [1.0,2.0,14.386294994851129,2.0,14.3862949948
51129]) |0.0 |
|red  |bad |16     | 14.386294994851129| (10, [0,2,3,4,7], [1.0,16.0,14.386294994851129,16.0,14.38629499
4851129]) |0.0 |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+
only showing top 10 rows
```

```
In [24]: # COMMAND -----  
  
## split train/test  
  
train, test = preparedDF.randomSplit([0.7, 0.3])
```

```
In [25]: # COMMAND -----  
  
## create model  
  
from pyspark.ml.classification import LogisticRegression  
lr = LogisticRegression(labelCol="label", featuresCol="features")
```

In [26]: `# COMMAND -----`

```
print (lr.explainParams())
```

aggregationDepth: suggested depth for treeAggregate ( $\geq 2$ ). (default: 2)  
 elasticNetParam: the ElasticNet mixing parameter, in range [0, 1]. For  $\alpha = 0$ , the penalty is an L2 penalty. For  $\alpha = 1$ , it is an L1 penalty. (default: 0.0)  
 family: The name of family which is a description of the label distribution to be used in the model. Supported options: auto, binomial, multinomial (default: auto)  
 featuresCol: features column name. (default: features, current: features)  
 fitIntercept: whether to fit an intercept term. (default: True)  
 labelCol: label column name. (default: label, current: label)  
 lowerBoundsOnCoefficients: The lower bounds on coefficients if fitting under bound constrained optimization. The bound matrix must be compatible with the shape (1, number of features) for binomial regression, or (number of classes, number of features) for multinomial regression. (undefined)  
 lowerBoundsOnIntercepts: The lower bounds on intercepts if fitting under bound constrained optimization. The bounds vector size must be equal with 1 for binomial regression, or the number of classes for multinomial regression. (undefined)  
 maxIter: max number of iterations ( $\geq 0$ ). (default: 100)  
 predictionCol: prediction column name. (default: prediction)  
 probabilityCol: Column name for predicted class conditional probabilities. Note: Not all models output well-calibrated probability estimates! These probabilities should be treated as confidences, not precise probabilities. (default: probability)  
 rawPredictionCol: raw prediction (a.k.a. confidence) column name. (default: rawPrediction)  
 regParam: regularization parameter ( $\geq 0$ ). (default: 0.0)  
 standardization: whether to standardize the training features before fitting the model. (default: True)  
 threshold: Threshold in binary classification prediction, in range [0, 1]. If threshold and thresholds are both set, they must match.e.g. if threshold is p, then thresholds must be equal to [1-p, p]. (default: 0.5)  
 thresholds: Thresholds in multi-class classification to adjust the probability of predicting each class. Array must have length equal to the number of classes, with values  $> 0$ , excepting that at most one value may be 0. The class with largest value p/t is predicted, where p is the original probability of that class and t is the class's threshold. (undefined)  
 tol: the convergence tolerance for iterative algorithms ( $\geq 0$ ). (default: 1e-06)  
 upperBoundsOnCoefficients: The upper bounds on coefficients if fitting under bound constrained optimization. The bound matrix must be compatible with the shape (1, number of features) for binomial regression, or (number of classes, number of features) for multinomial regression. (undefined)  
 upperBoundsOnIntercepts: The upper bounds on intercepts if fitting under bound constrained optimization. The bound vector size must be equal with 1 for binomial regression, or the number of classes for multinomial regression. (undefined)  
 weightCol: weight column name. If this is not set or empty, we treat all instance weights as 1.0. (undefined)

```
In [27]: # COMMAND -----
        ## train model
        fittedLR = lr.fit(train)
```

```
In [28]: fittedLR
```

```
Out[28]: LogisticRegressionModel: uid = LogisticRegression_5b6ada8feb48, numClasses = 2, numFeatures = 10
```

```
In [29]: # COMMAND -----
        train, test = df.randomSplit([0.7, 0.3])
```

```
In [30]: df.show(3, False)

+-----+-----+-----+-----+
|color|lab |value1|value2          |
+-----+-----+-----+-----+
|green|good|1      |14.386294994851129|
|blue |bad |8      |14.386294994851129|
|blue |bad |12     |14.386294994851129|
+-----+-----+-----+-----+
only showing top 3 rows
```

```
In [31]: # COMMAND -----
        rForm = RFormula()
        lr = LogisticRegression().setLabelCol("label").setFeaturesCol("features")
```

```
In [32]: # COMMAND -----
        from pyspark.ml import Pipeline
        stages = [rForm, lr]
        pipeline = Pipeline().setStages(stages)
```

```
In [33]: # COMMAND -----  
  
from pyspark.ml.tuning import ParamGridBuilder  
params = ParamGridBuilder()\br/>    .addGrid(rForm.formula, [  
        "lab ~ . + color:value1",  
        "lab ~ . + color:value1 + color:value2"])\br/>    .addGrid(lr.elasticNetParam, [0.0, 0.5, 1.0])\br/>    .addGrid(lr.regParam, [0.1, 2.0])\br/>    .build()
```

```
In [34]: # COMMAND -----  
  
from pyspark.ml.evaluation import BinaryClassificationEvaluator  
evaluator = BinaryClassificationEvaluator()\br/>    .setMetricName("areaUnderROC")\br/>    .setRawPredictionCol("prediction")\br/>    .setLabelCol("label")
```

```
In [35]: # COMMAND -----  
  
from pyspark.ml.tuning import TrainValidationSplit  
tvs = TrainValidationSplit()\br/>    .setTrainRatio(0.75)\br/>    .setEstimatorParamMaps(params)\br/>    .setEstimator(pipeline)\br/>    .setEvaluator(evaluator)
```

```
In [36]: # COMMAND -----  
  
tvsFitted = tvs.fit(train)
```

```
In [37]: type(tvsFitted)
```

```
Out[37]: pyspark.ml.tuning.TrainValidationSplitModel
```

```
In [ ]:
```



