

```
In [1]: from pyspark.sql import SparkSession
import pyspark.sql.functions as F
from pyspark.sql.types import *

spark = SparkSession\
    .builder\
    .appName("chapter-28-ML-recommender")\
    .getOrCreate()

import os
SPARK_BOOK_DATA_PATH = os.environ['SPARK_BOOK_DATA_PATH']
```

```
In [2]: from pyspark.ml.recommendation import ALS
from pyspark.sql import Row
ratings = spark.read\
    .text(SPARK_BOOK_DATA_PATH + "/data/sample_movielens_ratings.txt")\
    .rdd.toDF()\
    .selectExpr("split(value, '::') as col")\
    .selectExpr(
        "cast(col[0] as int) as userId",
        "cast(col[1] as int) as movieId",
        "cast(col[2] as float) as rating",
        "cast(col[3] as long) as timestamp"
    )

training, test = ratings.randomSplit([0.8, 0.2])

als = ALS()\
    .setMaxIter(5)\
    .setRegParam(0.01)\
    .setUserCol("userId")\
    .setItemCol("movieId")\
    .setRatingCol("rating")

print (als.explainParams())
alsModel = als.fit(training)           # fit to train Model on training data
predictions = alsModel.transform(test) # transform to predict on test data
```

alpha: alpha for implicit preference (default: 1.0)

checkpointInterval: set checkpoint interval (≥ 1) or disable checkpoint (-1). E.g. 10 means that the cache will get checkpointed every 10 iterations. Note: this setting will be ignored if the checkpoint directory is not set in the SparkContext. (default: 10)

coldStartStrategy: strategy for dealing with unknown or new users/items at prediction time. This may be useful in cross-validation or production scenarios, for handling user/item ids the model has not seen in the training data. Supported values: 'nan', 'drop'. (default: nan)

finalStorageLevel: StorageLevel for ALS model factors. (default: MEMORY_AND_DISK)

implicitPrefs: whether to use implicit preference (default: False)

intermediateStorageLevel: StorageLevel for intermediate datasets. Cannot be 'NONE'. (default: MEMORY_AND_DISK)

itemCol: column name for item ids. Ids must be within the integer value range. (default: item, current: movieId)

maxIter: max number of iterations (≥ 0). (default: 10, current: 5)

nonnegative: whether to use nonnegative constraint for least squares (default: False)

numItemBlocks: number of item blocks (default: 10)

numUserBlocks: number of user blocks (default: 10)

predictionCol: prediction column name. (default: prediction)
 rank: rank of the factorization (default: 10)
 ratingCol: column name for ratings (default: rating, current: rating)
 regParam: regularization parameter (≥ 0). (default: 0.1, current: 0.01)
 seed: random seed. (default: -1517157561977538513)
 userCol: column name for user ids. Ids must be within the integer value range. (default: user, current: userId)

In [3]: `# COMMAND -----`

```
alsModel.recommendForAllUsers(10)\
  .selectExpr("userId", "explode(recommendations)").show()
```

```
+-----+-----+
|userId|      col|
+-----+-----+
|    28|[12, 4.854412]|
|    28|[81, 4.5314903]|
|    28|[ 2, 3.90854]|
|    28|[82, 3.842551]|
|    28|[23, 3.5598414]|
|    28|[76, 3.4648323]|
|    28|[62, 3.214014]|
|    28|[26, 2.959738]|
|    28|[57, 2.8760154]|
|    28|[70, 2.7861943]|
|    26|[83, 5.9210377]|
|    26|[33, 5.4730806]|
|    26|[19, 5.261099]|
|    26|[37, 5.150425]|
|    26|[24, 5.059302]|
|    26|[12, 5.0024195]|
|    26|[88, 4.989318]|
|    26|[22, 4.9702883]|
|    26|[94, 4.957751]|
|    26|[23, 4.8820114]|
+-----+-----+
only showing top 20 rows
```

```
In [4]: alsModel.recommendForAllItems(10)\
        .selectExpr("movieId", "explode(recommendations)").show()
```

```
+-----+-----+
|movieId|      col|
+-----+-----+
|      31| [21, 4.190279]|
|      31| [12, 3.6101668]|
|      31| [6, 3.0709796]|
|      31| [8, 3.0144005]|
|      31| [14, 3.0032687]|
|      31| [7, 2.8663979]|
|      31| [13, 2.4790258]|
|      31| [10, 2.4686954]|
|      31| [25, 2.211947]|
|      31| [4, 2.1117184]|
|      85| [8, 4.849661]|
|      85| [16, 4.708046]|
|      85| [7, 3.7650425]|
|      85| [6, 3.5269358]|
|      85| [1, 3.0007]|
|      85| [14, 2.8658288]|
|      85| [21, 2.746056]|
|      85| [19, 2.5981874]|
|      85| [20, 2.3025043]|
|      85| [10, 2.057417]|
+-----+-----+
only showing top 20 rows
```

```
In [5]: # COMMAND -----

from pyspark.ml.evaluation import RegressionEvaluator
evaluator = RegressionEvaluator()\
    .setMetricName("rmse")\
    .setLabelCol("rating")\
    .setPredictionCol("prediction")
rmse = evaluator.evaluate(predictions)
print("Root-mean-square error = %f" % rmse)
```

```
Root-mean-square error = 1.544025
```

```
In [6]: # COMMAND -----  
  
from pyspark.mllib.evaluation import RegressionMetrics  
regComparison = predictions.select("rating", "prediction")\  
    .rdd.map(lambda x: (x(0), x(1)))  
metrics = RegressionMetrics(regComparison)
```

```
In [7]: # COMMAND -----  
  
from pyspark.mllib.evaluation import RankingMetrics, RegressionMetrics  
from pyspark.sql.functions import col, expr  
perUserActual = predictions\  
    .where("rating > 2.5")\  
    .groupBy("userId")\  
    .agg(expr("collect_set(movieId) as movies"))
```

```
In [8]: # COMMAND -----  
  
perUserPredictions = predictions\  
    .orderBy(col("userId"), expr("prediction DESC"))\  
    .groupBy("userId")\  
    .agg(expr("collect_list(movieId) as movies"))
```

```
In [ ]: # COMMAND -----  
  
perUserActualvPred = perUserActual.join(perUserPredictions, ["userId"]).rdd\  
    .map(lambda row: (row[1], row[2][:15]))
```

```
In [12]: type(perUserActualvPred)
```

```
Out[12]: pyspark.rdd.PipelinedRDD
```

```
In [14]: perUserActualvPred.collect()
```

```
Out[14]: [[([19, 49, 89, 92], [65, 19, 44, 3, 59, 49, 13, 89, 92]),
([18, 36, 73], [61, 18, 73, 36, 91]),
([75, 55, 44], [31, 44, 42, 95, 52, 55, 75]),
([66, 64, 50], [83, 66, 44, 84, 63, 4, 15, 88, 60, 67, 50, 64, 82, 87]),
([30, 69, 22], [14, 44, 45, 90, 22, 30, 63, 69, 96]),
([72], [78, 87, 73, 98, 26, 85, 31, 72, 71]),
([43], [75, 44, 38, 33, 18, 16, 68, 22, 94, 56, 20, 43, 53]),
([90], [90, 49, 38, 15, 65, 60, 2, 21]),
([36, 8, 80], [8, 91, 89, 81, 46, 83, 36, 72, 80, 62]),
([51, 90], [1, 73, 90, 78, 25, 24, 51]),
([54], [26, 2, 4, 11, 45, 61, 33, 54, 6, 55]),
([87, 14], [98, 14, 87, 77]),
([46, 56, 55], [70, 46, 56, 55, 93, 29, 66, 72, 19, 34]),
([96], [54, 96, 79, 33, 83, 4, 71, 15, 56, 9]),
([30, 13, 68, 55, 73, 23], [28, 59, 23, 68, 61, 13, 0, 73, 30, 55]),
([42], [66, 90, 42, 43, 98, 37, 8, 58, 41]),
([30, 96, 68, 69], [84, 78, 31, 96, 10, 30, 69, 68]),
([49], [70, 42, 49]),
([96, 29], [55, 20, 38, 66, 46, 40, 45, 51, 96, 29]),
([66, 27, 13, 50, 23],
[62, 70, 13, 45, 50, 43, 77, 66, 89, 6, 11, 88, 27, 23, 59]),
([52, 53, 93, 3, 72, 47],
[48, 49, 36, 39, 93, 6, 83, 47, 53, 5, 3, 95, 94, 45, 72]),
([83, 4], [76, 35, 78, 4, 83]),
([9, 92], [98, 9, 17, 45, 95, 34, 59, 92, 29])]
```

```
In [15]: ranks = RankingMetrics(perUserActualvPred)

# COMMAND -----

ranks.meanAveragePrecision, ranks.precisionAt(5)

# COMMAND -----
```

```
Out[15]: (0.29808913308913304, 0.49565217391304356)
```

```
In [ ]:
```

